

Data Exploration

```
> summary(train)
```

v1		v2		v3		v4	
Min.	:0.0002389	Min.	:1.000	Min.	:30.09	Min.	:1.000
1st Qu.:	:0.2294261	1st Qu.:	:1.062	1st Qu.:	:35.10	1st Qu.:	:1.079
Median :	:0.4803411	Median :	:1.134	Median :	:40.09	Median :	:1.158
Mean :	:0.4882555	Mean :	:1.136	Mean :	:40.08	Mean :	:1.155
3rd Qu.:	:0.7439418	3rd Qu.:	:1.207	3rd Qu.:	:45.29	3rd Qu.:	:1.234
Max.	:0.9984908	Max.	:1.284	Max.	:49.98	Max.	:1.300

v5		v6		v7		v8	
Min.	:2.999	Min.	:0.0000597	Min.	:-3.0328	Min.	:1.483
1st Qu.:	:3.454	1st Qu.:	:0.0241761	1st Qu.:	:-2.5380	1st Qu.:	:1.496
Median :	:3.963	Median :	:0.0504044	Median :	:-2.0406	Median :	:1.500
Mean :	:3.976	Mean :	:0.0497798	Mean :	:-2.0242	Mean :	:1.500
3rd Qu.:	:4.486	3rd Qu.:	:0.0746565	3rd Qu.:	:-1.5148	3rd Qu.:	:1.504
Max.	:5.000	Max.	:0.0998852	Max.	:-0.9635	Max.	:1.516

v9		Y	
Min.	:-5.5649	Min.	: 464.5
1st Qu.:	:-0.4858	1st Qu.:	: 629.8
Median :	: 0.5061	Median :	: 820.4
Mean :	: 0.5013	Mean :	: 836.2
3rd Qu.:	: 1.4764	3rd Qu.:	:1044.4
Max.	: 5.5063	Max.	:1269.7

Given the summary of the training data, it's necessary to preprocess the data since it present observations at different scales. We perform standardisation to get mean 0 and std 1.

```
# pre-processing so each predictor contribute equally
train_standardize <- as.data.frame(scale(train[1:10]))
train_standardize["Y"] <- list(train$Y)
```

For the first item we are going to check the assumption of linear regression model:

- Linearity
- Independence
- Normality
- Homoscedasticity

We will follow forward selection

Parametric Approach, Linear regression

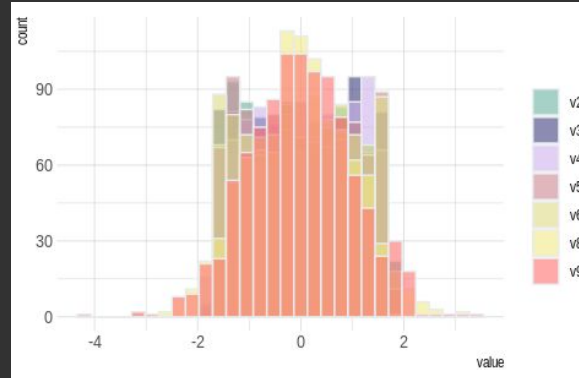
Independence - Correlation

```
# checking correlation between predictors  
corr <- cor(train[1:9])# just predictors  
which( corr>0.1 & corr<1, arr.ind=TRUE)
```

	row	col
v5	5	1
v7	7	1
v1	1	5
v7	7	5
v1	1	7
v5	5	7

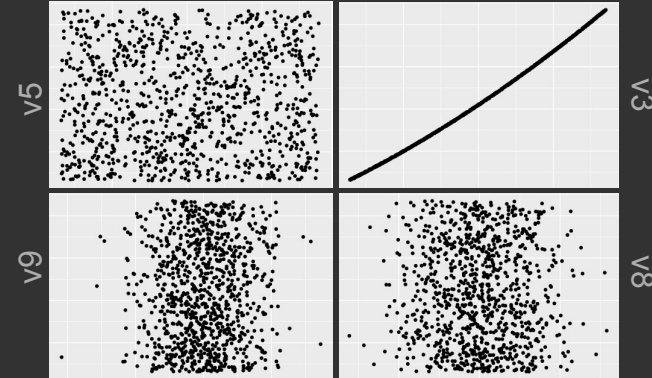
From the correlation analysis
v1, v7 and v5 are correlated.
We will keep the predictor with
significant p-value ($p < 0.05$).
The dropped predictors are:
V1 and v7

Normality



From the plot we can see some predictors
do not follow a normal distribution.
We will remove the predictors with
non-significant p-value and that do not
fulfil the condition above.
v3 do not follow it but have $p < 0.001$
The dropped predictors are:
v2, v4 and v6

Linearity

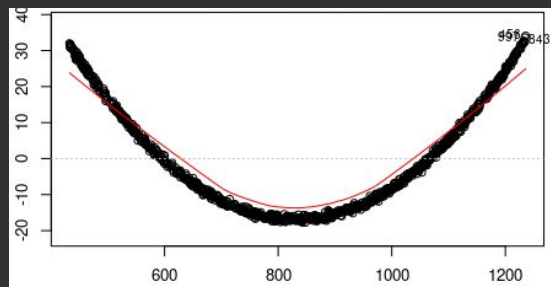


From the plot we can see some
predictors can not be described with a
straight line.
We will remove the predictors with
non-significant p-value that do not fulfil
the condition above
v5 do not fulfil but have $p < 0.05$
The dropped predictors are:
None

Parametric Approach, Linear regression

Homoscedasticity variance of error terms

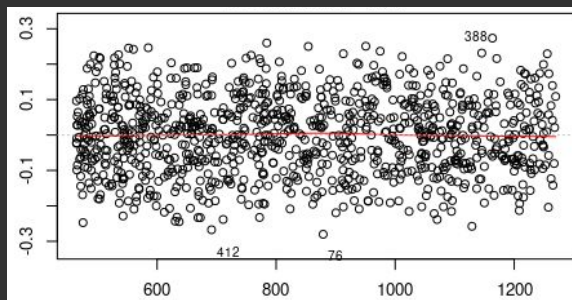
We try the model with the remaining predictors and plot Residual (y) vs Fitted values (x)



In the plot there is no visible randomness. A non-linear model would better describe the data. The relationship can be model as a polynomial regression

Polynomial regression

The residual describe a quadratic behavior. We try the 3 predictors as a quadratic term. The one with significant p-value ($p < 0.001$) is $v3$



Now there is more randomness in the residuals vs fitted, so the model fit more the data. The predictors fulfil the assumptions of the model

p-value significance

```
> linear_model <- lm(Y ~ v3 + I(v3^2) + v5 + v8 + v9, data=train_standardize)
> summary(linear_model)
```

```
Call:
lm(formula = Y ~ v3 + I(v3^2) + v5 + v8 + v9, data = train_standardize)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-0.280204 -0.078134  0.001416  0.080392  0.273492
```

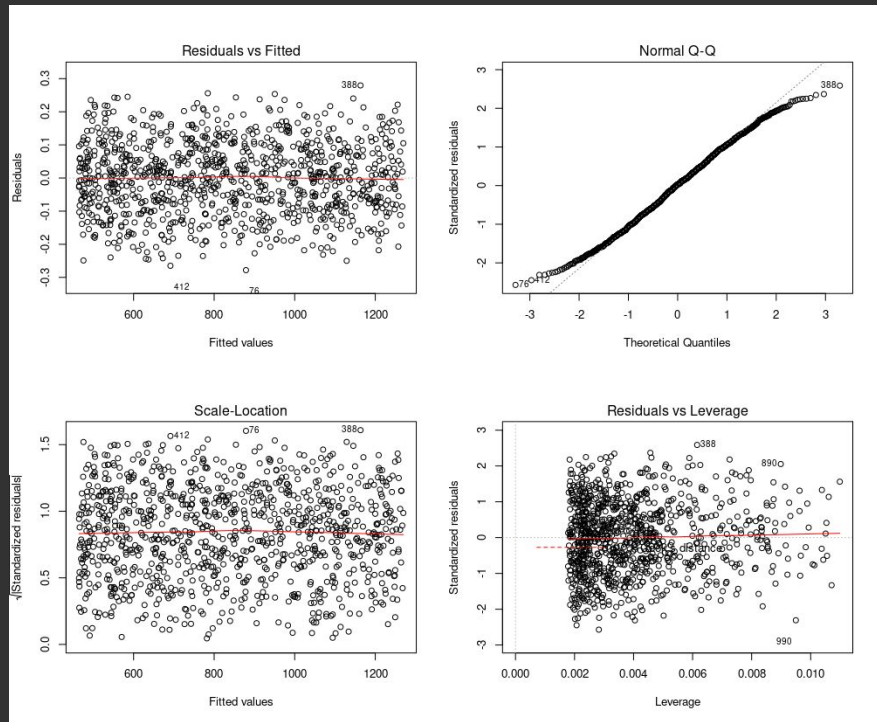
```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  819.334864   0.005193 157782.571 <2e-16 ***
v3           236.142109   0.003435  68737.139 <2e-16 ***
I(v3^2)      16.929288   0.003906  4333.674 <2e-16 ***
v5           1.187510   0.003430   346.231 <2e-16 ***
v8          -0.001020   0.003436   -0.297   0.767
v9           0.004983   0.003438    1.449   0.148
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.1083 on 994 degrees of freedom
Multiple R-squared:  1, Adjusted R-squared:  1
F-statistic: 9.477e+08 on 5 and 994 DF, p-value: < 2.2e-16
```

From the summary, $v8$ and $v9$ p-value ($p > 0.05$) are non-significant. Then we remove them. This predictors have no relationship with the prediction

Linear regression, Results

After the previous analysis the final linear regression model is:



```
> summary(linear_model)
```

Call:

```
lm(formula = Y ~ v3 + I(v3^2) + v5, data = train_standardize)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.278145	-0.077925	0.002741	0.078629	0.279523

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.193e+02	5.192e-03	157813.5	<2e-16 ***
v3	2.361e+02	3.432e-03	68799.6	<2e-16 ***
I(v3^2)	1.693e+01	3.905e-03	4335.6	<2e-16 ***
v5	1.187e+00	3.429e-03	346.3	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1083 on 996 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 1.579e+09 on 3 and 996 DF, p-value: < 2.2e-16

Note: Potential problems like outliers, high-leverage points were not present as eliminating v8 and v9 remove the need of the evaluation.

Non-parametric Approach, K-Nearest Neighbors

For the second item by definition no assumptions are made over the model. So we:

- Work in the curse of dimensionality by doing the combinations of predictors and evaluating the relevance of them.
- Find the value of k which reduce the MSE (mean square error)
- Do cross-validation with leave-one-out and majority voting, to find the best fitting model

This results were saved in the variable `compare`, the predictors, k and MSE values

```
# for each combination of predictors do the knn for values of k between 1 and sqrt(n)/2 = 15
# also as no test set is given cross validations leave-one-out of the
# training set is perform
for (combination in combinations_list){
  for (i in 1:length(k_values)){
    combination <- combination [! combination %in% zero]
    combination <- unlist(combination, use.names = FALSE)
    # run the knn
    knn_pred = knn.reg(train = train_standard[, combination], test = NULL, y = train$Y, k=k_values[i])
    # save the mse
    row <- data.frame(paste(unlist(combination), collapse=''), k_values[i], mean((knn_pred$residuals)^2))
    names(row) <- c("vars", "k", "mse")
    compare <- rbind(compare, row)
  }
}
```

```
> head(compare)
  vars k      mse
1   1 1 111516.83
2   1 2  86231.42
3   1 3  75428.60
4   1 4  70190.65
5   1 5  67718.01
6   1 6  66331.85
> compare[which(compare$mse == min(compare$mse)), ]
  vars k      mse
50   3 5  2.155141
```

The configuration with lowest MSE is using the v3 predictor and $k=5$

Conclusions

Linear regression

- Check the assumptions of the model.
- Be aware of the interpretation of summary tables and plots to realize potential problems, information loss.
- At the end 2 out of the 9 predictors were taken as having a relation with the label with a polynomial regression model.

K-Nearest neighbors

- Trying to avoid the curse of dimensionality just one predictor gave the best MSE. However the linear regression model concluded 2 predictors. This may be a situation of overfitting of the model.

Without the labels of the test set this cannot be verified.